

Ue01_Diamonds

November 14, 2022

```
[ ]: import pandas as pd
import math

df = pd.read_csv('diamonds.csv')
df.describe()
```

1 Implementation Part 1 (50%): Diamond Prices

1.1 1. Give an overview of the dataset structure by answering those questions:

1.1.1 How many samples and features are in the dataset?

```
[2]: print('Samples:', len(df))
print('Features:', len(df.columns))
```

Samples: 53940

Features: 10

1.1.2 What are the feature data types?

```
[3]: df.dtypes
```

```
[3]: carat      float64
cut          object
color        object
clarity       object
depth        float64
table        float64
price         int64
x            float64
y            float64
z            float64
dtype: object
```

1.1.3 Are diamonds balanced across color, cut and clarity? (Hint: roughly 1:1 means balanced, e.g. 1:2 is a “1:2 imbalance”)

```
[4]: color_df = df.groupby('color').size().to_frame('count')
color_df['balance'] = color_df['count'] / color_df['count'].min()
color_df = color_df.sort_values('balance')
print(color_df)
```

	count	balance
color		
J	2808	1.000000
I	5422	1.930912
D	6775	2.412749
H	8304	2.957265
F	9542	3.398148
E	9797	3.488960
G	11292	4.021368

Compared to the lowest count J, every class has at least an imbalance of 1:2 and most of them 1:3. E and F for example are balanced if viewed separately without the other colors.

```
[5]: cut_df = df.groupby('cut').size().to_frame('count')
cut_df['balance'] = cut_df['count'] / cut_df['count'].min()
cut_df = cut_df.sort_values('balance')
print(cut_df)
```

	count	balance
cut		
Fair	1610	1.000000
Good	4906	3.047205
Very Good	12082	7.504348
Premium	13791	8.565839
Ideal	21551	13.385714

Compared to the fair cut every class is imbalanced and ideal has an imbalance of 1:13

```
[6]: clarity_df = df.groupby('clarity').size().to_frame('count')
clarity_df['balance'] = clarity_df['count'] / clarity_df['count'].min()
clarity_df = clarity_df.sort_values('balance')
print(clarity_df)
```


	count	balance
clarity		
I1	741	1.000000
IF	1790	2.415655
VVS1	3655	4.932524
VVS2	5066	6.836707
VS1	8171	11.026991
SI2	9194	12.407557
VS2	12258	16.542510

SI1 13065 17.631579

The classes are extremely imbalanced, I1 vs VS2 has even an imbalance of 1:17.


1.2 2. Visualize diamond prices using a histogram, boxplot and density plot.

```
[7]: axarr = df.hist('price')
     for ax in axarr.flatten():
         ax.set_xlabel("price")
         ax.set_ylabel("count")
```




```
[8]: axarr = df.boxplot('price')
     axarr.set_ylabel("count")
```

```
[8]: Text(0, 0.5, 'count')
```



```
[9]: axarr = df['price'].plot.density()
     axarr.set_xlabel("price")
```

```
[9]: Text(0.5, 0, 'price')
```



1.2.1 Answer this question: Is there trend visible in those plots? If yes, which is it and in which plots can you see it?

There exist much more cheap diamonds up to 2500 and the amount decreases exponential which is very good visible in the boxplot and histogram. The boxplot also shows that there are many outliers beyond 12000. 75% of the diamonds are below 5000.

1.3 3. Calculate and state the mean, median, standard deviation, median absolute deviation (MAD), 1st and 3rd quartile (Q1 and Q3), and inner quartile range of the diamond price.

• If you are not familiar with those functions: use Google, Wikipedia, etc. • Required commands are all in the provided script.

```
[10]: print('Mean:', round(df['price'].mean(), 2))
      print('Median:', round(df['price'].median(), 2))
      print('STD: ', round(df['price'].std(), 2))
      # the mad() function is deprecated
      print('MAD: ', round((df['price'] - df['price'].mean()).abs().mean(), 2))
      print('Q1:', df['price'].quantile(0.25))
      print('Q3:', df['price'].quantile(0.75))
      print('Inner quartile: ', df['price'].quantile(0.75) - df['price'].quantile(0.
      ↪25))
```

Mean: 3932.8

Median: 2401.0

STD: 3989.44

MAD: 3031.6

Q1: 950.0

Q3: 5324.25


Inner quartile: 4374.25

1.4 4. Plot the diamond price against the carat values as a scatterplot. Answer this question:

Hint: plotting many samples will be slow. Changing the plot symbol to '?' will cause a speedup.

```
[11]: df.plot.scatter(x='carat', y='price', marker='.'))
```

```
[11]: <AxesSubplot:xlabel='carat', ylabel='price'>
```



1.4.1 Is there a trend visible in the plot? If yes, which is it?


There is no direct connection between price and carat. It depends on more than the carat feature when you look on the spread of the points. Tendencial more carat leads to higher prices.

1.5 5. Analyze the correlation between diamond price and diamond x, y, and z dimensions.

1.5.1 Create pairwise plots for these features.

```
[12]: pairwise_plot = df[['price', 'x', 'y', 'z']]
      pd.plotting.scatter_matrix(pairwise_plot, figsize=(15,15))
      print(pairwise_plot.corr())
```

	price	x	y	z
price	1.000000	0.884435	0.865421	0.861249
x	0.884435	1.000000	0.974701	0.970772
y	0.865421	0.974701	1.000000	0.952006
z	0.861249	0.970772	0.952006	1.000000



1.5.2 Is there a trend visible between x, y, and z? If yes, which is it?

Yes the dimensions all have a very high linear correlation, there is nearly a straight line.

1.5.3 Is there a trend visible between the dimensions and the price? If yes, which is it?

- Hint: if you don't know what a linear relation is (Google it!): – Linear correlation: feature A low \rightarrow feature B low, and feature A high \rightarrow feature B high. – (Inverse) linear correlation is also a linear correlation: feature A low \rightarrow feature B high, and feature A high \rightarrow feature B low: inverse linear correlation. Usually also just called linear correlation. – When plotting feature A against feature B and their points form a “straight line”, then it's a linear relationship between A and B = linear correlation.

The larger the dimensions the greater the price. X spreads more than y or z. There is also a large

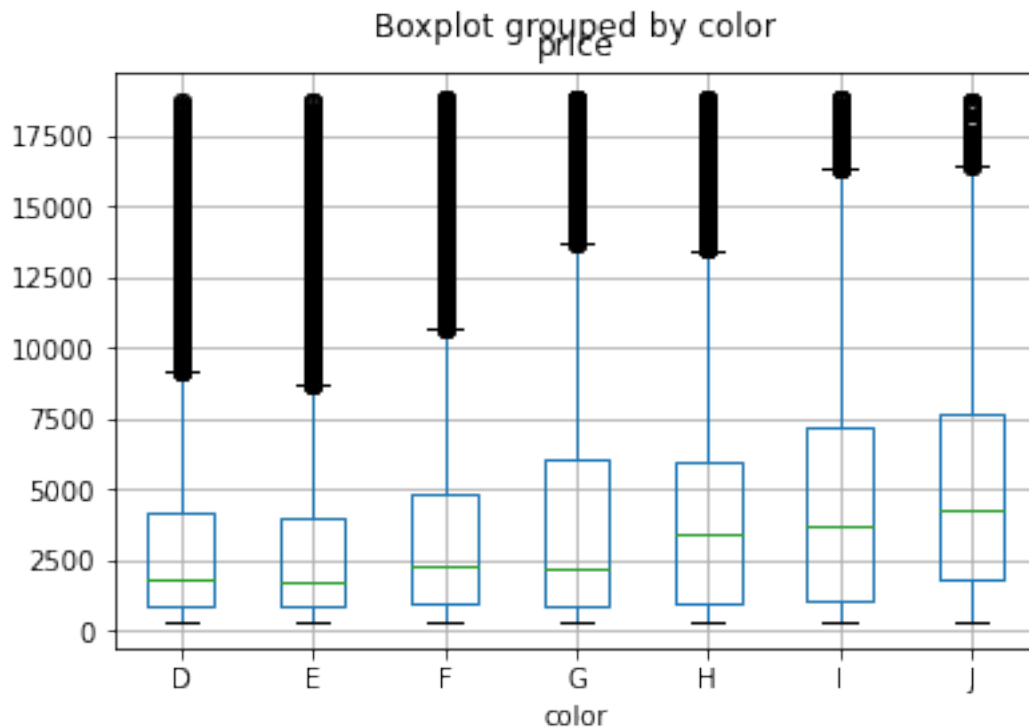
linear correlation between price and dimension, at least 86%.

1.6 6. Analyze diamond prices per diamond color.

1.6.1 Create boxplots showing diamond price boxes for each diamond color (all boxes should be in one figure).

```
[13]: df.boxplot('price', by='color')
```

```
[13]: <AxesSubplot:title={'center':'price'}, xlabel='color'>
```




1.6.2 Create densityplots showing diamond prices for each diamond color (all densities should be in one figure).

```
[14]: df.groupby('color')['price'].plot.density(legend=True)
```

```
[14]: color
D     AxesSubplot(0.125,0.125;0.775x0.755)
E     AxesSubplot(0.125,0.125;0.775x0.755)
F     AxesSubplot(0.125,0.125;0.775x0.755)
G     AxesSubplot(0.125,0.125;0.775x0.755)
H     AxesSubplot(0.125,0.125;0.775x0.755)
I     AxesSubplot(0.125,0.125;0.775x0.755)
```



```
J AxesSubplot(0.125,0.125;0.775x0.755)
Name: price, dtype: object
```



1.6.3 Answer this question: is there a trend visible? If yes, which one?

Most of the diamond colors are sold for approximately the same price. Depending on the color there exist more or less diamonds. Type E for example looks very stable and there are no peaks.

1.7 7. Use vectorized commands (= no loops!) to answer these questions:

1.7.1 How many diamonds have a price above 9500?

```
[15]: df[df['price'] > 9500].count()
```

```
[15]: carat      5734
      cut        5734
      color     5734
      clarity   5734
      depth     5734
      table     5734
      price     5734
      x         5734
      y         5734
      z         5734
      dtype: int64
```

1.7.2 How many diamonds have a price above 9500 and have color “D”?

```
[16]: df[(df['price'] > 9500) & (df['color'] == 'D')].count()
```

```
[16]: carat      461
      cut        461
      color     461
      clarity   461
      depth     461
      table     461
      price     461
      x         461
      y         461
      z         461
      dtype: int64
```

1.7.3 What is the mean and std of the price of all color “D” diamonds with cut “Fair”?

```
[17]: df[(df['cut'] == 'Fair') & (df['color'] == 'D')]['price'].mean()
```

```
[17]: 4291.061349693252
```

```
[18]: df[(df['cut'] == 'Fair') & (df['color'] == 'D')]['price'].std()
```

```
[18]: 3286.114238174996
```

1.7.4 What is the median and mad of the price of all color “J” diamonds with cut “Ideal”?

```
[19]: dfIdeal = df[(df['cut'] == 'Ideal') & (df['color'] == 'J')]['price']
      print('Median: ')
      print(dfIdeal.median())
      print('Ideal: ')
      print((dfIdeal - dfIdeal.mean()).abs().mean())
```

```
Median:
4096.0
Ideal:
3467.586682378051
```

1.7.5 Create two copies of the dataframe that contains only the price and carat feature. Apply a log with base 10 to both features in one of those dataframes, and square ($x' = x^2$) the features in the other dataframe. What is the mean and std of the transformed features in both dataframes?

```
[20]: dfSquare = df[['carat', 'price']].applymap(func=lambda x: x ** 2)
      print(dfSquare.mean())
      print(dfSquare.std())
```

```
carat    8.613903e-01
price    3.138225e+07
dtype: float64
carat    1.056506e+00
price    6.049189e+07
dtype: float64
```

```
[21]: dfLog = df[['carat', 'price']].applymap(func=math.log10)
      print(dfLog.mean())
      print(dfLog.std())
```

```
carat    -0.171532
price     3.381751
dtype: float64
carat     0.253987
price     0.440657
dtype: float64
```

```
[21]:
```

Ue01_CellBody

November 14, 2022

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

df = pd.read_csv('segmentationData.csv')
```

1 Implementation Part 2 (50%): Cell Body Segmentation Data.

1.1 1. Which classes exist? Are they (roughly) balanced?

```
[2]: df.describe()
```

```
[2]:
```

	AngleCh1	AreaCh1	AvgIntenCh1	AvgIntenCh2	AvgIntenCh3	\
count	2019.000000	2019.000000	2019.000000	2019.000000	2019.000000	
mean	90.493405	320.336305	126.071679	189.052115	96.420171	
std	48.760000	214.023533	165.008379	158.956105	96.666924	
min	0.030876	150.000000	15.160400	1.000000	0.120000	
25%	53.892205	193.000000	35.364158	44.998570	33.495693	
50%	90.588770	253.000000	62.343173	173.506300	67.431250	
75%	126.682013	362.500000	143.187800	279.289704	127.341651	
max	179.939323	2186.000000	1418.634831	989.509800	1205.512000	

	AvgIntenCh4	ConvexHullAreaRatioCh1	ConvexHullPerimRatioCh1	\
count	2019.000000	2019.000000	2019.000000	
mean	140.701585	1.205859	0.895764	
std	146.634665	0.202522	0.076108	
min	0.563265	1.005831	0.510623	
25%	40.679740	1.065236	0.856972	
50%	90.250000	1.148620	0.913262	
75%	191.170410	1.280514	0.955606	
max	886.837500	2.900320	0.996499	

	DiffIntenDensityCh1	FiberAlign2Ch3	IntenCoocMaxCh3	IntenCoocMaxCh4	\
count	2019.000000	2019.000000	2019.000000	2019.000000	
mean	72.660125	1.454076	0.231957	0.246709	
std	49.028338	0.252347	0.204030	0.183398	
min	25.760355	1.000000	0.014286	0.013423	

25%	43.532759	1.290022	0.051171	0.107596
50%	55.810304	1.469231	0.179775	0.211886
75%	79.909902	1.647809	0.353311	0.337116
max	442.773196	2.000000	0.968326	0.940367

	NeighborMinDistCh1	SkewIntenCh4
count	2019.000000	2019.000000
mean	29.691933	0.932515
std	11.501550	0.885901
min	10.083350	-1.004442
25%	22.547068	0.403460
50%	27.642860	0.728311
75%	34.079173	1.225431
max	126.993700	8.069013

```
[3]: df.dtypes
```

```
[3]: Class          object
AngleCh1          float64
AreaCh1            int64
AvgIntenCh1        float64
AvgIntenCh2        float64
AvgIntenCh3        float64
AvgIntenCh4        float64
ConvexHullAreaRatioCh1 float64
ConvexHullPerimRatioCh1 float64
DiffIntenDensityCh1 float64
FiberAlign2Ch3      float64
IntenCoocMaxCh3      float64
IntenCoocMaxCh4      float64
NeighborMinDistCh1  float64
SkewIntenCh4        float64
dtype: object
```

```
[4]: color_df = df.groupby('Class').size().to_frame('count')
color_df['balance'] = color_df['count'] / color_df['count'].min()
color_df = color_df.sort_values('balance')
print(color_df)
```

	count	balance
Class		
WS	719	1.000000
PS	1300	1.808067

The two classes WS and PS are not really balanced (imbalanced) with a ration of 1 : 1.8

1.2 2. Which noteworthy trends of features and relations between features as well as features and Class do you see?

```
[5]: df.loc[:, df.columns!='Class'].corr()
```

```
[5]:
```

	AngleCh1	AreaCh1	AvgIntenCh1	AvgIntenCh2	\
AngleCh1	1.000000	-0.025281	-0.026470	0.022270	
AreaCh1	-0.025281	1.000000	-0.039965	-0.163522	
AvgIntenCh1	-0.026470	-0.039965	1.000000	0.516892	
AvgIntenCh2	0.022270	-0.163522	0.516892	1.000000	
AvgIntenCh3	-0.008911	-0.139592	0.276232	0.191390	
AvgIntenCh4	0.006931	-0.084072	0.394118	0.599178	
ConvexHullAreaRatioCh1	-0.039384	0.320712	-0.238587	-0.448929	
ConvexHullPerimRatioCh1	0.032881	-0.489944	0.315972	0.438276	
DiffIntenDensityCh1	-0.040770	-0.074990	0.942705	0.485847	
FiberAlign2Ch3	-0.011776	-0.143195	-0.054437	-0.066706	
IntenCoocMaxCh3	0.005589	0.038674	0.124660	0.226547	
IntenCoocMaxCh4	0.009855	-0.015260	-0.176615	-0.242853	
NeighborMinDistCh1	0.054098	0.231252	-0.043701	-0.095988	
SkewIntenCh4	0.028063	0.075870	-0.141429	-0.236903	

	AvgIntenCh3	AvgIntenCh4	ConvexHullAreaRatioCh1	\
AngleCh1	-0.008911	0.006931	-0.039384	
AreaCh1	-0.139592	-0.084072	0.320712	
AvgIntenCh1	0.276232	0.394118	-0.238587	
AvgIntenCh2	0.191390	0.599178	-0.448929	
AvgIntenCh3	1.000000	0.390760	0.007011	
AvgIntenCh4	0.390760	1.000000	-0.259174	
ConvexHullAreaRatioCh1	0.007011	-0.259174	1.000000	
ConvexHullPerimRatioCh1	0.089375	0.274304	-0.716921	
DiffIntenDensityCh1	0.441698	0.386716	-0.193268	
FiberAlign2Ch3	-0.020619	-0.062946	0.050550	
IntenCoocMaxCh3	-0.326386	-0.165345	-0.283276	
IntenCoocMaxCh4	-0.134638	-0.477101	0.216181	
NeighborMinDistCh1	0.022199	0.024899	0.103121	
SkewIntenCh4	-0.079956	-0.420889	0.274447	

	ConvexHullPerimRatioCh1	DiffIntenDensityCh1	\
AngleCh1	0.032881	-0.040770	
AreaCh1	-0.489944	-0.074990	
AvgIntenCh1	0.315972	0.942705	
AvgIntenCh2	0.438276	0.485847	
AvgIntenCh3	0.089375	0.441698	
AvgIntenCh4	0.274304	0.386716	
ConvexHullAreaRatioCh1	-0.716921	-0.193268	
ConvexHullPerimRatioCh1	1.000000	0.276235	
DiffIntenDensityCh1	0.276235	1.000000	

FiberAlign2Ch3	0.027547	-0.046196
IntenCoocMaxCh3	0.216211	0.087012
IntenCoocMaxCh4	-0.118948	-0.160157
NeighborMinDistCh1	-0.163567	-0.051174
SkewIntenCh4	-0.169147	-0.114899


	FiberAlign2Ch3	IntenCoocMaxCh3	IntenCoocMaxCh4 \
AngleCh1	-0.011776	0.005589	0.009855
AreaCh1	-0.143195	0.038674	-0.015260
AvgIntenCh1	-0.054437	0.124660	-0.176615
AvgIntenCh2	-0.066706	0.226547	-0.242853
AvgIntenCh3	-0.020619	-0.326386	-0.134638
AvgIntenCh4	-0.062946	-0.165345	-0.477101
ConvexHullAreaRatioCh1	0.050550	-0.283276	0.216181
ConvexHullPerimRatioCh1	0.027547	0.216211	-0.118948
DiffIntenDensityCh1	-0.046196	0.087012	-0.160157
FiberAlign2Ch3	1.000000	-0.019200	0.107769
IntenCoocMaxCh3	-0.019200	1.000000	0.239047
IntenCoocMaxCh4	0.107769	0.239047	1.000000
NeighborMinDistCh1	-0.003185	-0.121486	-0.070880
SkewIntenCh4	0.060125	0.096885	0.724893

	NeighborMinDistCh1	SkewIntenCh4
AngleCh1	0.054098	0.028063
AreaCh1	0.231252	0.075870
AvgIntenCh1	-0.043701	-0.141429
AvgIntenCh2	-0.095988	-0.236903
AvgIntenCh3	0.022199	-0.079956
AvgIntenCh4	0.024899	-0.420889
ConvexHullAreaRatioCh1	0.103121	0.274447
ConvexHullPerimRatioCh1	-0.163567	-0.169147
DiffIntenDensityCh1	-0.051174	-0.114899
FiberAlign2Ch3	-0.003185	0.060125
IntenCoocMaxCh3	-0.121486	0.096885
IntenCoocMaxCh4	-0.070880	0.724893
NeighborMinDistCh1	1.000000	0.037793
SkewIntenCh4	0.037793	1.000000

For example AvgIntenCh1 and DiffIntenDensityCh1 seem to have a high correlation.

```
[6]: # https://stackoverflow.com/questions/29432629/
      ↪ plot-correlation-matrix-using-pandas
f = plt.figure(figsize=(10, 10))
plt.matshow(df.loc[:, df.columns!='Class'].corr(), fignum=f.number)
plt.xticks(range(df.select_dtypes(['number']).shape[1]), df.
      ↪ select_dtypes(['number']).columns, fontsize=12, rotation=45)
```

```
plt.yticks(range(df.select_dtypes(['number']).shape[1]), df.
    ↪select_dtypes(['number']).columns, fontsize=12)
cb = plt.colorbar()
cb.ax.tick_params(labelsize=14)
plt.title('Correlation Matrix', fontsize=16);
```



```
[7]: from sklearn import preprocessing

min_max_scaler = preprocessing.MinMaxScaler()
normalized_df = df.copy()
```



```

normalized_df[["AngleCh1", "AreaCh1", "AvgIntenCh1", "AvgIntenCh2",
↪ "AvgIntenCh3", "AvgIntenCh4", "ConvexHullAreaRatioCh1",
↪ "ConvexHullPerimRatioCh1", "DiffIntenDensityCh1", "FiberAlign2Ch3",
↪ "IntenCoocMaxCh3", "IntenCoocMaxCh4", "NeighborMinDistCh1", "SkewIntenCh4"]].
↪ = min_max_scaler.fit_transform(df[["AngleCh1", "AreaCh1", "AvgIntenCh1",
↪ "AvgIntenCh2", "AvgIntenCh3", "AvgIntenCh4", "ConvexHullAreaRatioCh1",
↪ "ConvexHullPerimRatioCh1", "DiffIntenDensityCh1", "FiberAlign2Ch3",
↪ "IntenCoocMaxCh3", "IntenCoocMaxCh4", "NeighborMinDistCh1", "SkewIntenCh4"]])


normalized_df.plot.box(by="Class", figsize=(15,10), logy=True, layout=(2,7))

```

```

[7]: AngleCh1                      AxesSubplot(0.125,0.53;0.0945122x0.35)
AreaCh1                          AxesSubplot(0.238415,0.53;0.0945122x0.35)
AvgIntenCh1                      AxesSubplot(0.351829,0.53;0.0945122x0.35)
AvgIntenCh2                      AxesSubplot(0.465244,0.53;0.0945122x0.35)
AvgIntenCh3                      AxesSubplot(0.578659,0.53;0.0945122x0.35)
AvgIntenCh4                      AxesSubplot(0.692073,0.53;0.0945122x0.35)
ConvexHullAreaRatioCh1          AxesSubplot(0.805488,0.53;0.0945122x0.35)
ConvexHullPerimRatioCh1        AxesSubplot(0.125,0.11;0.0945122x0.35)
DiffIntenDensityCh1            AxesSubplot(0.238415,0.11;0.0945122x0.35)
FiberAlign2Ch3                  AxesSubplot(0.351829,0.11;0.0945122x0.35)
IntenCoocMaxCh3                 AxesSubplot(0.465244,0.11;0.0945122x0.35)
IntenCoocMaxCh4                 AxesSubplot(0.578659,0.11;0.0945122x0.35)
NeighborMinDistCh1             AxesSubplot(0.692073,0.11;0.0945122x0.35)
SkewIntenCh4                    AxesSubplot(0.805488,0.11;0.0945122x0.35)
dtype: object

```



```
[8]: df.groupby("Class").mean()
```

```
[8]:
```

	AngleCh1	AreaCh1	AvgIntenCh1	AvgIntenCh2	AvgIntenCh3	\
Class						
PS	90.619486	314.339231	78.342220	138.852531	96.309678	
WS	90.265441	331.179416	212.369728	279.816315	96.619950	

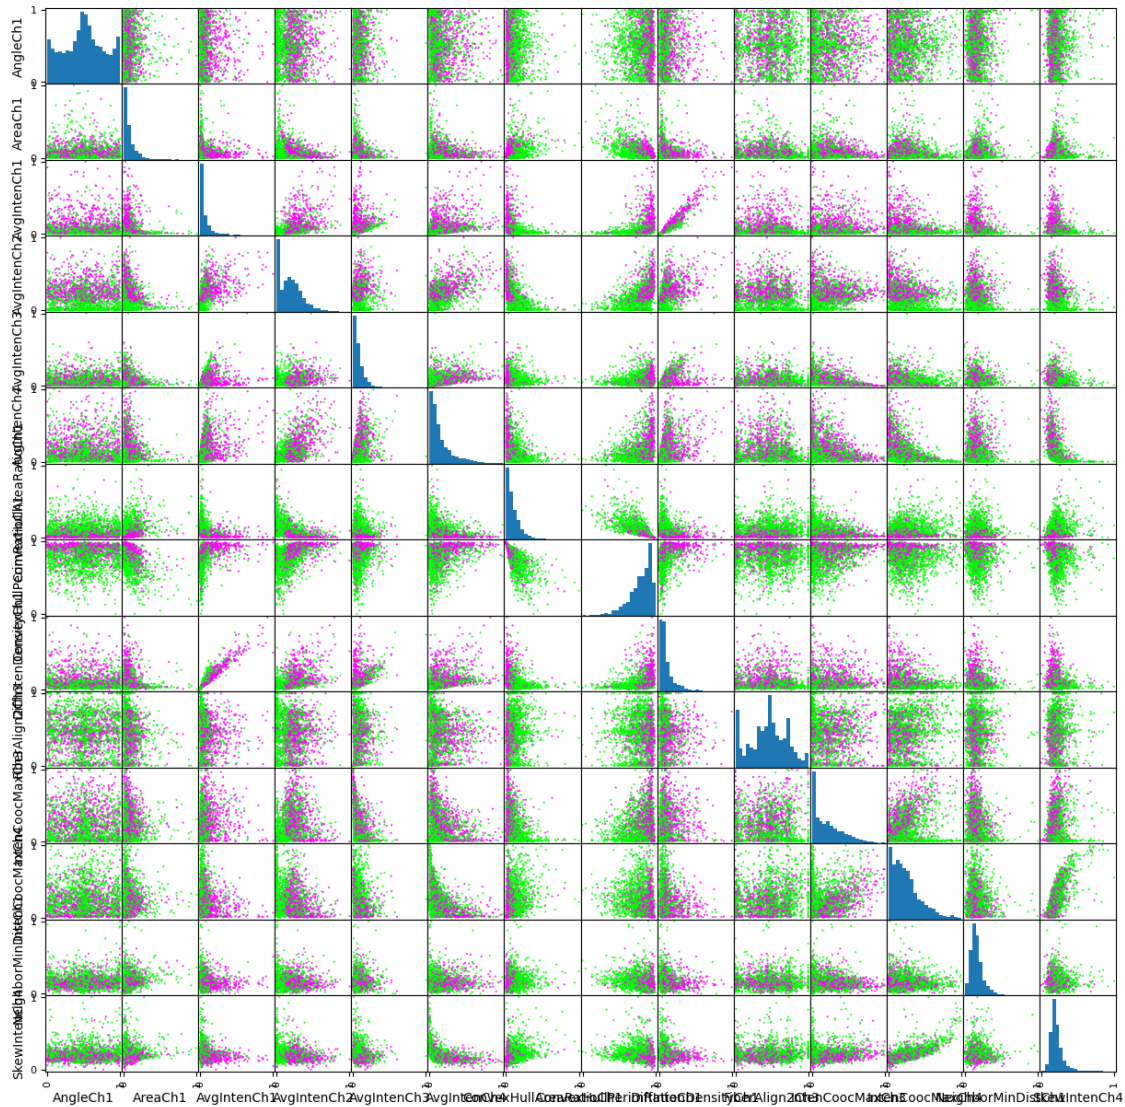
	AvgIntenCh4	ConvexHullAreaRatioCh1	ConvexHullPerimRatioCh1	\
Class				
PS	114.191873	1.255922	0.875728	
WS	188.632915	1.115343	0.931988	

	DiffIntenDensityCh1	FiberAlign2Ch3	IntenCoocMaxCh3	IntenCoocMaxCh4	\
Class					
PS	60.748834	1.470619	0.189937	0.269191	
WS	94.196536	1.424166	0.307932	0.206060	

	NeighborMinDistCh1	SkewIntenCh4
Class		
PS	30.063084	1.054383
WS	29.020866	0.712169

```
[9]: columns = ["AngleCh1", "AreaCh1", "AvgIntenCh1", "AvgIntenCh2", "AvgIntenCh3",
    ↪ "AvgIntenCh4", "ConvexHullAreaRatioCh1", "ConvexHullPerimRatioCh1",
    ↪ "DiffIntenDensityCh1", "FiberAlign2Ch3", "IntenCoocMaxCh3",
    ↪ "IntenCoocMaxCh4", "NeighborMinDistCh1", "SkewIntenCh4"]

test = lambda x: '#0f0' if x == "PS" else "#f0f"
vfunc = np.vectorize(test)
grr = pd.plotting.scatter_matrix(
    normalized_df[columns], c=vfunc(df.Class), figsize=(15,15), marker='.',
    hist_kwds={'bins':20}, s=10, alpha=.8)
```



```
[10]: print('Mean:', round(df['AreaCh1'].mean(), 2))
print('Median:', round(df['AreaCh1'].median(), 2))
print('STD: ', round(df['AreaCh1'].std(), 2))
# the mad() function is deprecated
print('MAD: ', round((df['AreaCh1'] - df['AreaCh1'].mean()).abs().mean(), 2))
print('Q1:', df['AreaCh1'].quantile(0.25))
print('Q3:', df['AreaCh1'].quantile(0.75))
print('Inner quartile: ', df['AreaCh1'].quantile(0.75) - df['AreaCh1'].
      ↪ quantile(0.25))
```

```
Mean: 320.34
Median: 253.0
STD: 214.02
```

MAD: 137.58

Q1: 193.0

Q3: 362.5

Inner quartile: 169.5

Some features like Angle1, AvgIntenCh3 or NeighborMinDistCh1 don't help to distinguish between the two classes PS and WS. AvgIntenCh2 for example has a different median and a different value range which could work very well to separate the two classes.

1.3 3. If you would need to distinguish the classes with those features, which features would you choose, and why?

- AvgIntenCh1
- AvgIntenCh2
- DiffIntenDensityCh1
- ConvexHullAreaRatioCh1
- ConvexHullPerimRatioCh1

When you look at the median and inner quartile range of the boxplot for these features, the classes are easily separable.

[]: